

Scala - A Powerful and Scalable Function-Objective Programming Language

Troy Hu and Benjamin Killeen

Abstract—Insert abstract here.

I. INTRODUCTION

Software components are simply self-contained parts of software used by larger parts or entire applications. Modern software typically relies on many common components like hashable types or iterable data structures to use as building blocks. Component abstraction enables the generalization of implementation for these features, greatly reducing duplication of effort. It is one of the most powerful tools in the programmer’s utility belt, a fact which the designers of the Scala programming language aimed to exploit [1]. In the early 2000s, existing languages had little support for type-sound component abstraction, including widely used languages like Java and C#. Odersky *et al.* address this shortcoming with a language that combines elements of object-oriented and functional programming in a statically typed environment. Scala, which gets its name from “scalable,” provides a powerful interface for abstraction within a development framework intended for mass adoption.

In pursuit of usability, Scala borrows many syntactic elements from Java and C#, and it integrates smoothly with components from these languages. In fact, the Scala library includes standard Java objects like `java.lang.String`, as shown in Listing 1. Scala code can take advantage of existing implementations in Java, and in the end it compiles to Java bytecode, making Scala packages available to Java programmers. At the same time, Scala maintains a distinct programming paradigm from either Java or C#. It discards some features from these languages, and it develops completely novel ideas from the ν Obj calculus [2]. The example in Listing 1 highlights syntactic similarities between Java and Scala, comparing two implementations of the same program. Note how Java prepends type declarations before terms, whereas Scala affixes type declarations using the `:` operator. This and other changes effect a terser, more expressive syntax overall. Of course, Scala’s foremost strength comes from its typing system. Abstract class definitions and path-dependent types utilize the ν Obj Calculus, enabling incredible flexibility through the use of traits and mixins [2]. Somewhat akin to Java’s abstract classes, traits allow a programmer to rely on abstract methods for common functionalities. For example, the `Equiv[T]` trait in Fig. 3 represents an equivalence relation on the type `T`, abstracting the definition of `eq` on which a concrete method, `neq` depends. Mixins enable a class to inherit from multiple traits. For instance, one might use

```
// Java:
class PrintOptions {
  public static void main(String[] args) {
    System.out.println("Options selected:");
    for (int i = 0; i < args.length; i++)
      if (args[i].startsWith("-"))
        System.out.println(" " +
          ↪ args.substring(1))
  }
}

// Scala:
object PrintOptions {
  def main(args: Array[String]): unit = {
    println("Options selected:")
    for (val arg <- args)
      if (arg.startsWith("-"))
        println(" " + arg.substring(1))
  }
}
```

Listing 1: Notice how Scala’s general syntax and structure are similar to Java’s. At the same time, there are some visible differences, e.g., `unit` is returned in the Scala implementation instead of `void` in the Java implementation.

`Equiv` in conjunction with an `Ordering` trait to represent separate relations on the same type in one object.

Finally, the uniform object model in Scala provides a cohesive programming environment. Every Scala value is an object, and every operation is a call to a method. The boolean `true`, for example, is a singleton object that extends (inherits from) the `Boolean` trait (see Sec. ??). At the same time, Scala fits into a functional programming paradigm. Functions themselves are values, and the syntax allows for SML-like decomposition and pattern matching. These features result in a powerful language with a unified programming experience.

II. RELATED WORK

The ideas underlying Scala come from a variety of work. First and foremost, the designers rely on the ν Obj calculus outlined in Odersky *et al.* (2003) for a theoretical grounding on which they build Scala’s robust type system [2]. Although a thorough discussion of the ν Obj calculus is beyond the scope of this paper, which focuses on the Scala language proper, we describe its main points in brief. Additionally, the ideas in Scala have been replicated or expanded on in more recent work, which we discuss here.

The ν Obj calculus is, broadly, a calculus and dependent type system that includes classes and objects with types as mem-

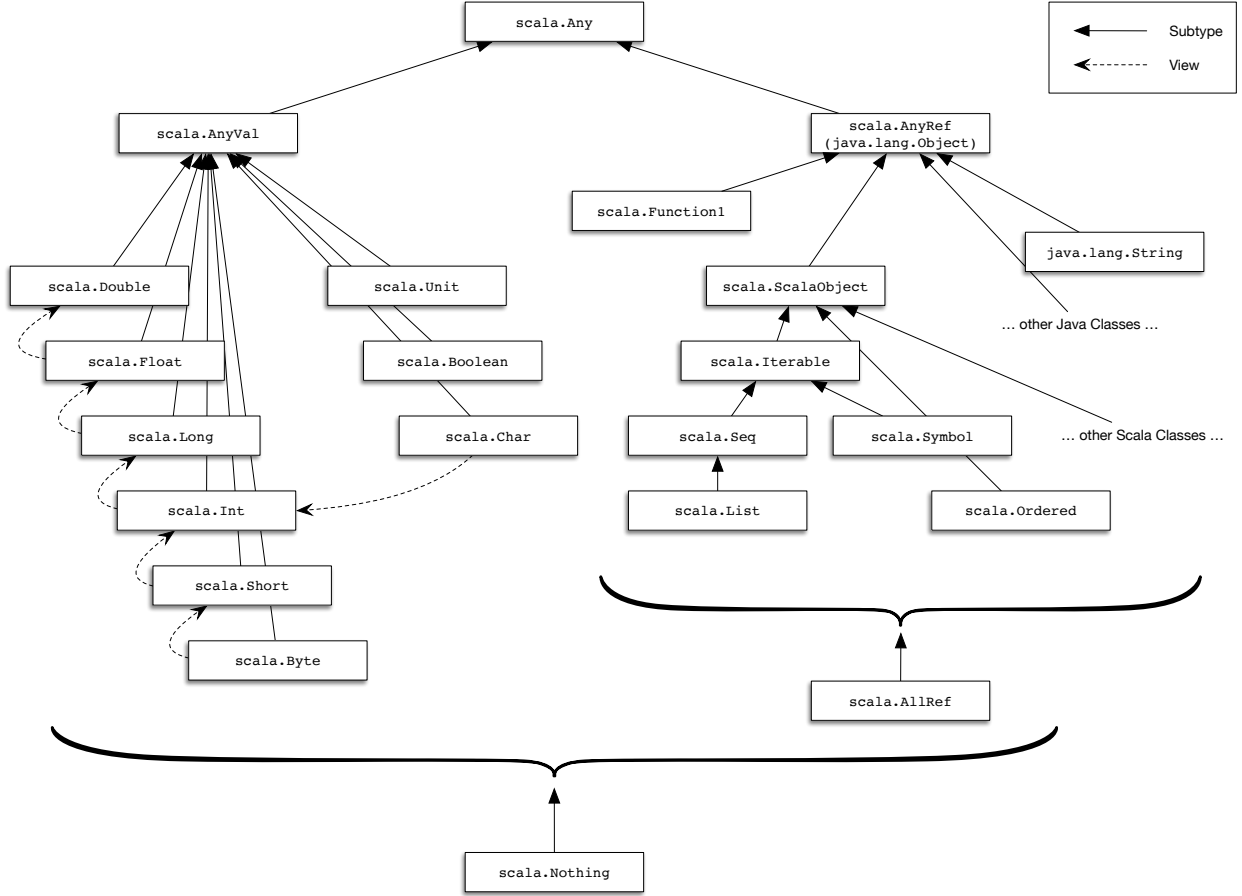


Fig. 1: A visual representation of Scala's class hierarchy

bers [2]. *Dependent types* rely on some value for definition, and with this formalism, the νObj calculus expresses Java's inner class system, virtual types, and family polymorphism. Furthermore, it models SML-style modules and functors, diverging from standard object-oriented type systems in three fundamental ways:

- 1) Objects *and* classes are considered primitive, as opposed to just objects.
- 2) Type checking and evaluation rely on name references rather than object records.
- 3) Object types are expressible using possibly nominal type components.

These alterations make possible the rigorous blend of object-oriented and functional styles evident in Scala, differing markedly from traditional languages like Java and C#.

The Rust programming language is a popular alternative to Scala that is designed to target systems-level applications. The language is similar to C++ in that it allows users to interact directly with hardware. Moreover, the syntax of Rust is very similar to that of C++. Rust, however, is unlike C++ in that any program purely written in Rust is guaranteed to avoid memory errors and data races. Most notably, Rust utilizes many of Scala's core concepts to achieve its goals. Like Scala, Rust is both a functional and object-oriented programming language (to model C++). Rust also ensures type soundness in order to avoid memory errors such as dangling pointers. Unlike Scala,

however, Rust emphasizes concurrency and parallelism. The remainder of this paper focuses on the practical considerations of programming in the Scala language.

III. PROGRAMMING IN SCALA

Being a fully developed programming language, Scala implements numerous useful features that encourage modularity and reusability in software. In this section, we highlight just a few of these, as well as provide an overview of the general process through which software construction takes place in Scala.

A. Unified Object Model

In Scala, every value is an object and every class is a subtype of the `Any` class. Fig. 1 shows an overview of the unified object model in Scala's standard library, under which many classes simply absorb the existing Java implementation. Immutable objects or *values*, denoted with the `val` keyword on declaration, inherit from the class `AnyVal`, and include classes like `scala.Boolean` and `scala.Char`. Mutable *variables* on the other hand, inherit from `scala.AnyRef`, which is synonymous with the generic `java.lang.Object`. These include classes like `scala.List` or traits like `scala.Ordered`. Even functions fit into this object model, such as the `scala.Function1` trait, which single-argument functions implement.

One of the fundamental aspects of Scala is its invariance with respect to value representation. Under this rule, *interpreting an value belonging to some subclass as an element of its superclass does not change that value's representation*. Put another way, this invariance rule demands that for types S , T with $S <: T$, and x of type S ,

```
x.asInstanceOf[T].asInstanceOf[S] = x
```

always holds. Because of this rule, Scala's object model includes some unusual relations not found in more traditional languages. In particular, Fig. 1 uses dashed lines to denote a *View* from one type to another. As discussed in Sec. III-C, a view enables interpretation between types which are not strictly subtypes of one another. For example, both `Int` and `Float` maintain a member called `MaxValue`, which other function can rely on without requiring a strict subtype of either `Int` or `Float`. Finally, the `scala.Nothing` trait is a subtype of every other class, allowing for flexible objects like the `scala.List` object `Nil`, of type `List[Nothing]`.

Another fundamental feature of Scala's object model inclusion of operations, which are invocations of a method. For example, the usual addition operator `+` can be implemented like any other method, as in Lst. 5, and invoked in the exact same way: `x.+(y)`. Of course Scala allows for the more conventional calling style `x + y`, but it does so merely through syntactic sugaring. Here, x is the receiver object, $+$ is a method defined in x , and y is the method's argument.

Finally, Scala differs from Java with regard to the construction of objects. Rather than define a dedicated constructor function within objects, the class name itself is simply the constructor. Object instantiations simply runs the entire body of the class, instantiated member variables, values, and definitions.

B. Traits, Objects, and Classes

Traits in Scala are somewhat akin to abstract classes in Java, defining abstract members that should be implemented by more concrete classes. As a very simple example, consider the `Nat` trait which defines necessary operations on natural numbers:

```
trait Nat {
  def isZero: Boolean;
  def pred: Nat;
  def succ: Nat;
}
```

Classes which inherit from `Nat` must implement the `isZero`, `succ`, and `pred` methods with the corresponding types in order to inherit from `Nat`. This is advantageous for several reasons. First, as we shall see, the `Nat` class may use these abstract methods to define more complex functions, which subclasses need not implement. Second, other functions may refer to `Nats` as a whole rather than either subclass of the original trait.

Objects or singleton objects in Scala often accompany trait definitions. A singleton object is at the same time a class definition and the sole instantiated member of that class. For instance, the `Z` object represents the natural number “zero,” of which there can only be one:

```
object Z extends Nat {
  def isZero = true
  def succ = S(this)
  def pred = this
}
```

This simple implementation behaves as we would expect, relying on the `S` class (below) to implement `Z.succ`. The choice to ensure `Z.pred = Z` was made in keeping with conventional evaluation rules. The use of the `extends` keyword represents the inheritance from `Nat`.

Similarly classes in Scala can “extend” a trait, while taking in arguments to the constructor. The `S` class

```
case class S(n: Nat) extends Nat {
  def isZero = false
  def succ = S(this)
  def pred = n
}
```

represents the “successor” of another natural n , either `Z` or another `S`. Here, the use of the `case` keyword results in a constructor very similar to constructors in SML, which can be instantiated without the use of `new` and easily deconstructed in pattern matching blocks.

C. Views

In place of loose subtyping, Scala requires views to implicitly convert objects from one type to another. A view is implemented with a method that takes in an argument of one type and returns an object of another type. The only difference between a view method and a normal method is that view methods require the `implicit` modifier, which goes before the method definition. This modifier allows the Scala compiler to know that it is the implicit conversion method when converting from one type to another. Scala implicitly applies a view to an expression, e of type T , when one of the following cases occur:

- The expected type of e is not of type T .
- A member selected from e is not a member of T .

For example, in Lst. 2, `listToSet` is the view that converts `GenList[T]` to `Set[T]`. The compiler inserts applications of the view onto `xs`.

D. Pattern Matching in an Object-Oriented Setting

Unlike Java and other object-oriented programming languages, Scala implements pattern matching. That is, Scala provides the programmer with a natural and functional-like mechanism for “creating structured data representations similar to algebraic data types and a decomposition mechanism based on pattern matching.”

Since Scala is an object-oriented programming language, it does not have algebraic data types. Instead, Scala creates structured data representations through the `case` modifier. If `case` precedes the definition of a class, a factory method with the same arguments as the primary class constructor is

```

trait Set[T] {
  def include(x: T): Set[T]
  def contains(x: T): boolean
}

implicit def listToSet[T](xs: GenList[T]):
  → Set[T] =
  new Set[T] {
    def include(x: T): Set[T] =
      xs prepend x
    def contains(x: T): boolean =
      !isEmpty && (xs.head == x || (xs.tail
        → contains x))
  }

```

Listing 2: An example of views and implicit conversions.

automatically defined. For example, in Figure 4, since the **Num** and **Plus** classes are defined with the **case** modifier, we can define an anonymous **Num** object without using the **new** keyword. As the reader can see, factory methods are very similar in structure to the constructors of algebraic data types. In fact, factory methods serve the same purpose as constructors when pattern matching.

Scala’s pattern matching expressions can decompose the factory method constructors as patterns. The syntax for pattern matching expressions is

```

x match {
  case p1 => e1
  case p2 => e2
  ...
}

```

where p_1, p_2 are patterns compared against x in order, and e_1, e_2 are the resulting expressions. Each p_i is of the form `FactoryMethod(x_1, x_2, \dots, x_n)`, where `FactoryMethod` refers to the factory method constructors discussed previously. When a match p_i is found, then e_i is executed. For example, in Figure 5, the `eval` function matches `term` against `Num(x)` and `Plus(left, right)` (the constructors are from Figure 4).

```

abstract class Term
case class Num(x: int) extends Term
case class Plus(left: Term, right: Term) extends Term

```

Fig. 2: example

```

object Interpreter {
  def eval(term: Term): int = term match {
    case Num(x) => x
    case Plus(left, right) => eval(left) + eval(right)
  }
}

```

Fig. 3

Note: unlike Java, writing a simple language interpreter in Scala is almost as easy as writing an interpreter in SML. Examples of pattern matching being used for an interpreter can be seen in the calculator program we wrote.

```

trait Equiv[T] {
  def eq (x: T, y: T): Boolean; // abstract
  def neq (x: T, y: T): Boolean = !eq(x, y)
}

```

Listing 3: A simple equivalence relation Scala.

```

trait Ordering[T] {
  def leq (x: T, y: T): Boolean; // abstract
  → method
  def geq (x: T, y: T): Boolean = leq(y, x)
  def gt (x: T, y: T): Boolean = !leq(x, y)
  def lt (x: T, y: T): Boolean = !leq(y, x)
  def eq (x: T, y: T): Boolean = leq(x, y)
  → && leq(y, x)
  def neq (x: T, y: T): Boolean = !leq(x, y)
  → || !leq(y, x)
}

```

Listing 4: An ordering relation in Scala. It is important to distinguish between an ordered type and an ordering on that type, of which there can be arbitrarily many. `Ordering[T]` represents the latter.

IV. ENGAGEMENT: PEANO ARITHMETIC CALCULATOR

V. DISCUSSION

A. Current Status of Work

After 15 years, Scala has been regularly updated and is currently at stable release version 2.12.8. Over this period, the core principles of Scala have remained the same. Scala has achieved widespread adoption in the industry with companies such as Twitter and Apple utilizing the language. Moreover, Scala has both a large academic and non-academic user base. The language is often cited or used in computer science research. In addition, Scala’s user community is thriving, there exist chatrooms, subreddits, and research conferences devoted entirely to Scala. Numerous libraries, tutorials, and guides are run and maintained by the community. The main community page can be found at: <https://www.scala-lang.org/community/>. A detailed, updated, and easy to read documentation can be found at: <https://docs.scala-lang.org/>. Finally, there are dedicated installers/installation guides for all operating systems at: scala-lang.org/download.

REFERENCES

- [1] M. Odersky, P. Altherr, V. Cremet, B. Emir, S. Maneth, S. Micheloud, N. Mihaylov, M. Schinz, E. Stenman, and M. Zenger, “An overview of the Scala programming language,” Tech. Rep., 2004.
- [2] M. Odersky, V. Cremet, C. Röckl, and M. Zenger, “A Nominal Theory of Objects with Dependent Types,” in *ECOOP 2003 – Object-Oriented Programming*, ser. Lecture Notes in Computer Science, L. Cardelli, Ed. Springer Berlin Heidelberg, 2003, pp. 201–224.

```

package nat

import ord._

trait Nat {
  def isZero: Boolean;
  def pred: Nat;
  def succ: Nat;
  def + (that: Nat): Nat = if (that.isZero)
    ↪ this else this.succ + that.pred
  def - (that: Nat): Nat = if (that.isZero)
    ↪ this else this.pred - that.pred
  def * (that: Nat): Nat = if (that.isZero)
    ↪ that else this + this * that.pred
  def ^ (that: Nat): Nat = if (that.isZero)
    ↪ that.succ else this * (this ^
    ↪ that.pred)
  def <= (that: Nat): Boolean =
    ↪ NatOrdering.leq(this, that)
  def >= (that: Nat): Boolean =
    ↪ NatOrdering.geq(this, that)
  def < (that: Nat): Boolean =
    ↪ NatOrdering.lt(this, that)
  def > (that: Nat): Boolean =
    ↪ NatOrdering.gt(this, that)
  def == (that: Nat): Boolean =
    ↪ NatOrdering.eq(this, that)
  def != (that: Nat): Boolean =
    ↪ NatOrdering.neq(this, that)
}

object NatOrdering extends Ordering[Nat] {
  def leq (x: Nat, y: Nat): Boolean =
    (x.isZero, y.isZero) match {
      case (true, _) => true
      case (false, true) => false
      case (false, false) => leq(x.pred,
        ↪ y.pred)
    }
}

object Zero extends Nat {
  override def toString () = "Z"
  def isZero = true
  def succ = Succ(this)
  def pred = this
}

case class Succ(n: Nat) extends Nat {
  override def toString () = s"S($n)"
  def isZero = false
  def succ = Succ(this)
  def pred = n
}

```

Listing 5: An example outlining Scala classes.