

Scala - A Powerful and Scalable Function-Objective Programming Language

Troy Hu and Benjamin Killeen

Abstract—Insert abstract here.

I. INTRODUCTION

Software components are simply self-contained parts of software used by larger parts or entire applications. Modern software typically relies on many common components like hashable types or iterable data structures to use as building blocks. Component abstraction enables the generalization of implementation for these features, greatly reducing duplication of effort. It is one of the most powerful tools in the programmer’s utility belt, a fact which the designers of the Scala programming language aimed to exploit [1]. In the early 2000s, existing languages had little support for type-sound component abstraction, including widely used languages like Java and C#. Odersky *et al.* address this shortcoming with a language that combines elements of object-oriented and functional programming in a statically typed environment. Scala, which gets its name from “scalable,” provides a powerful interface for abstraction within a development framework intended for mass adoption.

In pursuit of usability, Scala borrows many syntactic elements from Java and C#, and it integrates smoothly with components from these languages. In fact, the Scala library includes standard Java objects like `java.lang.String`, as shown in 1. Scala code can take advantage of existing implementations in Java, and in the end it compiles to Java bytecode, making Scala packages available to Java programmers. At the same time, Scala maintains a distinct programming paradigm from either Java or C#. It discards some features from these languages, and it develops completely novel ideas from the ν Obj calculus [2]. The example in Fig. 1 highlights syntactic similarities between Java and Scala, comparing two implementations of the same program. Note how Java prepends type declarations before terms, whereas Scala affixes type declarations using the `:` operator. This and other changes effect a terser, more expressive syntax overall.

Of course, Scala’s foremost strength comes from its typing system. Abstract class definitions and path-dependent types utilize the ν Obj Calculus, enabling incredible flexibility through the use of traits and mixins [2]. Somewhat akin to Java’s abstract classes, traits allow a programmer to rely on abstract methods for common functionalities. For example, the `Equiv[T]` trait in Fig. 2 represents an equivalence relation on the type `T`, abstracting the definition of `eq` on which a concrete method, `neq` depends. Mixins enable a class to inherit from multiple traits. For instance, one might use

```
// Java:
class PrintOptions {
  public static void main(String[] args) {
    System.out.println("Options selected:");
    for (int i = 0; i < args.length; i++)
      if (args[i].startsWith("-"))
        System.out.println(" " + args.substring(1))
  }
}

// Scala:
object PrintOptions {
  def main(args: Array[String]): unit = {
    println("Options selected:")
    for (val arg <- args)
      if (arg.startsWith("-"))
        println(" " + arg.substring(1))
  }
}
```

Listing 1: Notice how Scala’s general syntax and structure are similar to Java’s. At the same time, there are some visible differences, e.g., `unit` is returned in the Scala implementation instead of `void` in the Java implementation.

`Equiv` in conjunction with an `Ordering` trait to represent separate relations on the same type in one object.

Finally, the uniform object model in Scala provides a cohesive programming environment. Every Scala value is an object, and every operation is a call to a method. The boolean `true`, for example, is a singleton object that extends (inherits from) the `Boolean` trait (see Sec. III-A). At the same time, Scala fits into a functional programming paradigm. Functions themselves are values, and the syntax allows for SML-like decomposition and pattern matching. These features result in a powerful language with a unified programming experience.

II. RELATED WORK

The ideas underlying Scala come from a variety of work. First and foremost, the designers rely on the ν Obj calculus outlined in Odersky *et al.* (2003) for a theoretical grounding on which they build Scala’s robust type system [2]. Although a thorough discussion of the ν Obj calculus is beyond the scope of this paper, which focuses on the Scala language proper, we describe its main points in brief. Additionally, the ideas in Scala have been replicated or expanded on in more recent work, which we discuss here.

The ν Obj calculus is, broadly, a calculus and dependent type system that includes classes and objects with types as members [2]. *Dependent types* rely on some value for definition, and with this formalism, the ν Obj calculus expresses Java’s inner class system, virtual types, and family polymorphism.

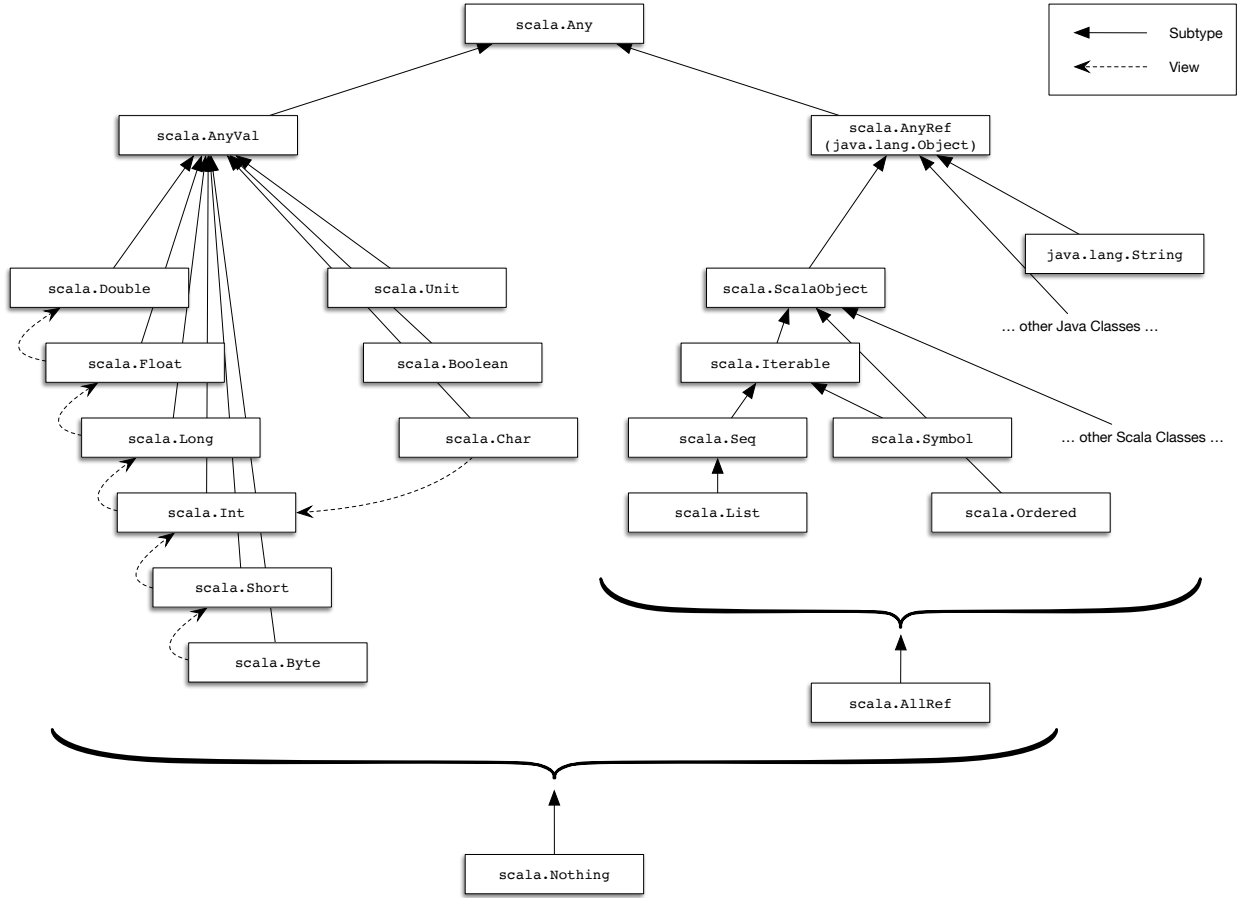


Fig. 1: A visual representation of Scala's class hierarchy

Furthermore, it models SML-style modules and functors, diverging from standard object-oriented type systems in three fundamental ways:

- 1) Objects *and* classes are considered primitive, as opposed to just objects.
- 2) Type checking and evaluation rely on name references rather than object records.
- 3) Object types are expressible using possibly nominal type components.

These alterations make possible the rigorous blend of object-oriented and functional styles evident in Scala, differing markedly from traditional languages like Java and C#.

The remainder of this paper focuses on the practical considerations of programming in the Scala language.

III. PROGRAMMING IN SCALA

A. Unified Object Model

In Scala, every value is an object and every class is a subtype of the `Any` class. Below the `Any` class, Scala classes can generally be divided into two groups: value classes that inherit from the `AnyVal` class and reference classes that inherit from the `AnyRef` class. Note that the value classes that are not `AnyVal` do not subtype each other. This means, for example, that the `Int` class is not a subtype of the `Float` class. Scala prevents such subtyping because the language

forces an invariant that when a value is interpreted in a subclass and a super class, the values's representation does not change. For example, both `Int` and `Float` maintain a `MaxValue` member. However, the `MaxValue` of `Float` is different from the `MaxValue` of `Int`. Instead, Scala implements **views** (discussed below) which allow from implicit conversions from one type to another. As expected of an object-oriented language, class `Null` is a subtype of every reference type. However, since Scala is part functional, the language also has a `Nothing` type, which represents an empty type, is a subtype of every single type. For example, `Nil` is defined as `List[Nothing]`. Note that `Nothing` allows for the implementation of options, another functional concept, in Scala.

Another major aspect of Scala's object model is that every operation is the invocation of a method. For example, `x + y` is syntactic sugar for `x.+(y)`; `x` is the receiver object, `+` is a method defined in `x`, and `y` is the method's argument. In fact, Scala desugars every identifier between two expressions as a method call of the first expression. Unlike constructors in Java, Scala constructors just come from the class name. There is no need for a separate class constructor inside the class definition. When a class is instantiated with its constructor, the whole body of the class is executed.

B. Singleton Objects

Scala is very unique in that it implements Singleton Objects. Singleton Objects are classes that can only have one instance while a Scala program is running. Such classes are defined almost exactly the same as normal Scala classes. The only difference is that the `class` modifier is replaced by the `object` modifier. Some examples of singleton objects include: a class representation of Zero and the `PrintOptions` example from above. Note that Singleton Objects are used extensively when Scala uses a Java class. Every static member of a Java class is stored in a Singleton Object.

```
package nat
import ord._

trait Nat {
  def isZero: Boolean;
  def pred: Nat;
  def succ: Nat;
  def + (that: Nat): Nat = if (that.isZero) this else
    → this.succ + that.pred
  def - (that: Nat): Nat = if (that.isZero) this else
    → this.pred - that.pred
  def * (that: Nat): Nat = if (that.isZero) that else this
    → + this * that.pred
  def ^ (that: Nat): Nat = if (that.isZero) that.succ else
    → this * (this ^ that.pred)
  def <= (that: Nat): Boolean = NatOrdering.leq(this,
    → that)
  def >= (that: Nat): Boolean = NatOrdering.geq(this,
    → that)
  def < (that: Nat): Boolean = NatOrdering.lt(this, that)
  def > (that: Nat): Boolean = NatOrdering.gt(this, that)
  def == (that: Nat): Boolean = NatOrdering.eq(this, that)
  def != (that: Nat): Boolean = NatOrdering.neq(this,
    → that)
}

object NatOrdering extends Ordering[Nat] {
  def leq (x: Nat, y: Nat): Boolean =
    (x.isZero, y.isZero) match {
      case (true, _) => true
      case (false, true) => false
      case (false, false) => leq(x.pred, y.pred)
    }
}

object Zero extends Nat {
  override def toString () = "0"
  def isZero = true
  def succ = Succ(this)
  def pred = this
}

case class Succ(v: Nat) extends Nat {
  override def toString () = s"$v"
  def isZero = false
  def succ = Succ(this)
  def pred = v
}
```

Listing 2: An example outlining Scala classes.

C. Pattern Matching in an Object-Oriented Setting

Unlike Java and other object-oriented programming languages, Scala implements pattern matching. That is, Scala provides the programmer with a natural and functional-like mechanism for “creating structured data representations similar to algebraic data types and a decomposition mechanism based on pattern matching.”

Since Scala is an object-oriented programming language, it does not have algebraic data types. Instead, Scala creates

structured data representations through the `case` modifier. If `case` precedes the definition of a class, a factory method with the same arguments as the primary class constructor is automatically defined. For example, in Figure 4, since the `Num` and `Plus` classes are defined with the `case` modifier, we can define an anonymous `Num` object without using the `new` keyword. As the reader can see, factory methods are very similar in structure to the constructors of algebraic data types. In fact, factory methods serve the same purpose as constructors when pattern matching.

Scala’s pattern matching expressions can decompose the factory method constructors as patterns. The syntax for pattern matching expressions is

$$x \text{ case } \{ \text{case } p_1 \Rightarrow e_1 \text{ case } p_2 \Rightarrow e_2 \dots \}$$

This syntax matches the object x against the patterns p_1, p_2, \dots in order. Each p_i is of the form $FactoryMethod(x_1, x_2, \dots, x_n)$, where $FactoryMethod$ refers to the factory method constructors discussed previously. When a match p_i is found, then e_i is executed. For example, in Figure 5, the `eval` function matches term against `Num(x)` and `Plus(left, right)` (the constructors are from Figure 4).

```
abstract class Term
case class Num(x: int) extends Term
case class Plus(left: Term, right: Term) extends Term
```

Fig. 2: example

```
object Interpreter {
  def eval(term: Term): int = term match {
    case Num(x) => x
    case Plus(left, right) => eval(left) + eval(right)
  }
}
```

Fig. 3

Note: unlike Java, writing a simple language interpreter in Scala is almost as easy as writing an interpreter in SML. Examples of pattern matching being used for an interpreter can be seen in the calculator program we wrote.

D. Views

Views are used to implicitly convert one type to another. A view is implemented with a method that takes in an argument of one type and returns an object of another type. The only difference between a view method and a normal method is that view methods require the `implicit` modifier, which goes before the method definition. This modifier allows the Scala compiler to know that it is the implicit conversion method when converting from one type to another. Scala implicitly applies a view to an expression, e of type T , when one of the following cases occur:

- The expected type of e is not of type T .
- A member selected from e is not a member of T .

For example, in Figure 7, `listToSet` is the view that converts `GenList[T]` to `Set[T]`. The compiler inserts applications of the view onto `xs`.

```

package Set

trait Set[T]{
  def include(x: T): Set[T]
  def contains(x: T): boolean
}

implicit def listToSet[T](xs: GenList[T]): Set[T] =
  new Set[T] {
    def include(x: T): Set[T] =
      xs prepend x
    def contains (x: T): boolean =
      !isEmpty && (xs.head == x || (xs.tail contains x))
  }

/*****
//Assume that xs is a value of type GenList[T]
val s: Set[T] = xs;
xs contains x

//Compiler inserts applications of the view =>

val s: Set[T] = listToSet(xs);
listToSet(xs) contains x;
*****/

```

Fig. 4: An example of views and implicit conversions.

```

trait Equiv[T] {
  def eq (x: T, y: T): Boolean; // abstract
  def neq (x: T, y: T): Boolean = !eq(x, y)
}

```

Listing 3: A simple equivalence relation Scala.

IV. ENGAGEMENT: PEANO ARITHMETIC CALCULATOR

V. DISCUSSION

A. Current Status of Work

After 15 years, Scala has been regularly updated and is currently at stable release version 2.12.8. Over this period, the core principles of Scala have remained the same. Scala has achieved widespread adoption in the industry with companies such as Twitter and Apple utilizing the language. Moreover, Scala has both a large academic and non-academic user base. The language is often cited or used in computer science research. In addition, Scala’s user community is thriving, there exist chatrooms, subreddits, and research conferences devoted entirely to Scala. Numerous libraries, tutorials, and guides are run and maintained by the community. The main community page can be found at: <https://www.scala-lang.org/community/>. A detailed, updated, and easy to read documentation can be found at: <https://docs.scala-lang.org/>. Finally, there are dedicated installers/installation guides for all operating systems at: <https://www.scala-lang.org/download/>.

REFERENCES

- [1] M. Odersky, P. Altherr, V. Cremet, B. Emir, S. Maneth, S. Micheloud, N. Mihaylov, M. Schinz, E. Stenman, and M. Zenger, “An overview of the Scala programming language,” Tech. Rep., 2004.
- [2] M. Odersky, V. Cremet, C. Röckl, and M. Zenger, “A Nominal Theory of Objects with Dependent Types,” in *ECOOP 2003 – Object-Oriented Programming*, ser. Lecture Notes in Computer Science, L. Cardelli, Ed. Springer Berlin Heidelberg, 2003, pp. 201–224.

```

trait Ordering[T] {
  def leq (x: T, y: T): Boolean; // abstract method
  def geq (x: T, y: T): Boolean = leq(y, x)
  def gt (x: T, y: T): Boolean = !leq(x, y)
  def lt (x: T, y: T): Boolean = !leq(y, x)
  def eq (x: T, y: T): Boolean = leq(x, y) && leq(y, x)
  def neq (x: T, y: T): Boolean = !leq(x, y) || !leq(y, x)
}

```

Listing 4: An ordering relation in Scala. It is important to distinguish between an ordered type and an ordering on that type, of which there can be arbitrarily many. `Ordering[T]` represents the latter.